

```
In [1]: import pandas as pd
```

```
In [2]: df=pd.read_csv('C:/shubhangi/2023-24/LP-III_ML/Assignment 2/emails.csv')
```

```
In [3]: df.head()
```

Out[3]:

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastructure	military
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	0	0	0	0	0
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	0	0	0	0	0
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	0	0	0	0	0
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	0	0	0	0	0
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	0	0	0	0	0

5 rows × 3002 columns



```
In [4]: df.shape
```

Out[4]: (5172, 3002)

```
In [5]: #input data
x=df.drop(['Email No.','Prediction'],axis=1)
#output data
y=df['Prediction']
```

```
In [6]: x.shape
```

Out[6]: (5172, 3000)

```
In [7]: x.dtypes
```


```
Out[7]: the          int64
to          int64
ect         int64
and         int64
for         int64
...
infrastructure int64
military      int64
allowing      int64
ff            int64
dry           int64
Length: 3000, dtype: object
```

```
In [8]: set(x.dtypes)
```

```
Out[8]: {dtype('int64')}
```

```
In [10]: import seaborn as sns  
sns.countplot(x=y)
```

```
Out[10]: <Axes: xlabel='Prediction', ylabel='count'>
```



```
In [11]: y.value_counts()
```

```
Out[11]: 0    3672  
         1    1500  
         Name: Prediction, dtype: int64
```

```
In [14]: pip install scikit-learn
```

```
Requirement already satisfied: scikit-learn in c:\programdata\anaconda3\lib\site-packa  
ges (1.2.2)  
Requirement already satisfied: numpy>=1.17.3 in c:\programdata\anaconda3\lib\site-pack  
ages (from scikit-learn) (1.24.3)  
Requirement already satisfied: scipy>=1.3.2 in c:\programdata\anaconda3\lib\site-packa  
ges (from scikit-learn) (1.10.1)  
Requirement already satisfied: joblib>=1.1.1 in c:\programdata\anaconda3\lib\site-pack  
ages (from scikit-learn) (1.2.0)  
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\si  
te-packages (from scikit-learn) (2.2.0)  
Note: you may need to restart the kernel to use updated packages.
```

```
In [15]: from sklearn.preprocessing import MinMaxScaler
```

```
In [16]: scaler=MinMaxScaler()
```

```
In [17]: x_scaled=scaler.fit_transform(x)
```

```
In [18]: x_scaled
```

```
Out[18]: array([[0.         , 0.         , 0.         , ..., 0.         , 0.         ,
                0.         ],
               [0.03809524, 0.09848485, 0.06705539, ..., 0.         , 0.00877193,
                0.         ],
               [0.         , 0.         , 0.         , ..., 0.         , 0.         ,
                0.         ],
               ...,
               [0.         , 0.         , 0.         , ..., 0.         , 0.         ,
                0.         ],
               [0.00952381, 0.0530303 , 0.         , ..., 0.         , 0.00877193,
                0.         ],
               [0.1047619 , 0.18181818, 0.01166181, ..., 0.         , 0.         ,
                0.         ]])
```

```
In [19]: #cross validation
         from sklearn.model_selection import train_test_split
```

```
In [20]: x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,random_state=0,test_size=0.25
```

```
In [22]: x_scaled.shape
```

```
Out[22]: (5172, 3000)
```

```
In [23]: x_train.shape
```

```
Out[23]: (3879, 3000)
```

```
In [24]: x_test.shape
```

```
Out[24]: (1293, 3000)
```

```
In [26]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [27]: knn=KNeighborsClassifier(n_neighbors=5)
```

```
In [28]: knn.fit(x_train,y_train)
```

```
Out[28]: 

▼ KNeighborsClassifier


         KNeighborsClassifier()
```


```
In [29]: #predict on test data
         y_pred=knn.predict(x_test)
```

```
In [31]: #import evaluation matrix
         from sklearn.metrics import ConfusionMatrixDisplay,accuracy_score
```

```
In [32]: from sklearn.metrics import classification_report
```

```
In [33]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[33]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x207c27b6c50>
```



```
In [34]: y_test.value_counts()
```

```
Out[34]: 0    929
         1    364
         Name: Prediction, dtype: int64
```

```
In [35]: accuracy_score(y_test,y_pred)
```

```
Out[35]: 0.871616395978345
```

```
In [36]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.84	0.90	929
1	0.70	0.95	0.81	364
accuracy			0.87	1293
macro avg	0.84	0.89	0.85	1293
weighted avg	0.90	0.87	0.88	1293

```
In [37]: import numpy as np
import matplotlib.pyplot as plt
```

```
In [40]: error = []
         for k in range(1,41):
             knn= KNeighborsClassifier(n_neighbors=k)
             knn.fit(x_train,y_train)
             pred=knn.predict(x_test)
             error.append(np.mean(pred!=y_test))
```

```
In [39]: error
```

```
Out[39]: [0.10827532869296211,
          0.10982211910286156,
          0.12296983758700696,
          0.11523588553750967,
          0.12838360402165508,
          0.1214230471771075,
          0.15158546017014696,
          0.14849187935034802,
          0.17246713070378963,
          0.16705336426914152,
          0.1871616395978345,
          0.18329466357308585,
          0.21500386697602475,
          0.21345707656612528,
          0.22815158546017014,
          0.2266047950502707,
          0.23588553750966745,
          0.23356535189481825,
          0.2459396751740139,
          0.24361948955916474,
          0.2559938128383604,
          0.2552204176334107,
          0.2699149265274555,
          0.2691415313225058,
          0.2822892498066512,
          0.28306264501160094,
          0.2954369682907966,
          0.2923433874709977,
          0.3039443155452436,
          0.300077339520495,
          0.30549110595514306,
          0.30549110595514306,
          0.31245166279969067,
          0.31245166279969067,
          0.3194122196442382,
          0.317092034029389,
          0.32637277648878577,
          0.32559938128383603,
          0.33410672853828305,
          0.3325599381283836]
```

```
In [41]: knn=KNeighborsClassifier(n_neighbors=1)
```

```
In [42]: knn.fit(x_train,y_train)
```

```
Out[42]: KNeighborsClassifier
          KNeighborsClassifier(n_neighbors=1)
```

```
In [43]: y_pred=knn.predict(x_test)
```

```
In [44]: accuracy_score(y_test,y_pred)
```

```
Out[44]: 0.8917246713070379
```

```
In [45]: from sklearn.svm import SVC
```

```
In [*]: svm=SVC(kernel='linear')
```

```
In [*]: svm.fit(x_train,y_train)
```

```
In [*]: y_pred=svm.predict(x_test)
```

```
In [54]: accuracy_score(y_test,y_pred)
```

```
Out[54]: 0.9767981438515081
```

```
In [ ]: #Linear 0.9767981438515081
```